

Paper

http://www.vpri.org/_pdf/tr2011001_final_worlds.pdf

[paper reviewed on Oct. 14, 2021 by Toronto CS Cabal. The following are my notes on this paper.]

Why Am I Attracted to this Paper?

Isolation

Objects in a world have no implicit dependencies on objects in other worlds.

PEG Makes This Easy

The ideas in this paper seep into many parts of a GPL (General Purpose programming Language).

PEG - OMeta (ancestor of Ohm-JS) - makes it easy to rewrite the syntax of any GPL, hence, making it easy to enact these changes.

Generates Ideas

I think that I originally saw this in the Ohm-JS thesis.

This shows what is possible when you elide details and "free your mind" to think about higher-level concepts.

Worlds is not directly related to Ohm-JS, but springs forth as an idea of how to use Ohm-JS to build new kinds of solutions to problems.

Overview of Paper

Abstract

1 Introduction

2 Approach

3 Worlds by Example

4 Property Lookup Semantics

5 Implementation

6 Case Study - Bitmap Editor

7 Case Study - OMeta

8 Related Work

9 Conclusions, Future

10 Acknowledgements

References

1 Introduction

"... An important class of problems have to perform *speculations* and *experiments*, often in parallel, to discover how to proceed. ..." [pt: *I call this Design, Brainstorming, Architecture, etc.*)]

paraphrase: Try/Catch is a subset of undo. [pt: Kind of like Greenspun's 10th Rule (for undo).]

Web surfing ... back button ... exploration ...

"... This is somewhat similar to *transactions* ..."

"... Worlds are first-class structures ..."

[pt: *I believe that every useful programming concept needs to be made explicit, e.g. GC (garbage collection), OO, etc. Making something first-class in a language is but one way to make concepts explicit. (GC is not usually first-class, OO objects are usually first-class)*]

2 Approach

Worlds realized in

- JavaScript
- Squeak (Smalltalk).
- *Sprout* a world (instantiate from prototype),
- make changes
- *commit* changes back to parent (if possible, see below for algorithm)
- field access
 - lookups
 - lookup in local scope
 - chain upwards through parents, if not found in local scope.
 - updates
 - always done locally
 - *commit* operation pushes changes to parent.
 - [pt: *Is this similar to pre-CL Lisps and/or special variables in CL?*
E.G. "dynamic scoping"]

2.1 JS

```
A = thisWorld;  
p = new Point (1, 2);  
  
B = A.sprout ();  
in B {  
  p.y = 3  
}  
  
C = A.sprout ();  
in C {  
  p.y = 7;  
}  
  
C.commit ();
```

P.y is 2 in A, while p.y is 3 in B, while p.y is 7 in C.

P.y in A becomes 7 after commit.

2.2 Safety Properties

- No Surprises
- Consistency

3 Worlds by Example

3.1 Better Support for Exceptions

3.2 Undo for Applications

3.3 Extension Methods in JavaScript

- scoped methods

4 Property Lookup Semantics

- turnstile notation

5 Implementation

5.1 Data Structures

- WObject
- WWorld

WObject

Each slot of each object contains 2 fields:

- Reads
- Writes

[pt: We are accustomed to thinking of variables as containing exactly 1 field, but WObjects contain 2 fields]

Each slot is characterized by

- don't know - '?'
- a value

WWorld

collection of WObjects

5.2 Slot Update

5.2 Update $w.x.i$


5.2.1 (optimization) - create *writes* for x

5.2.2 write v into $w.x.i.writes$

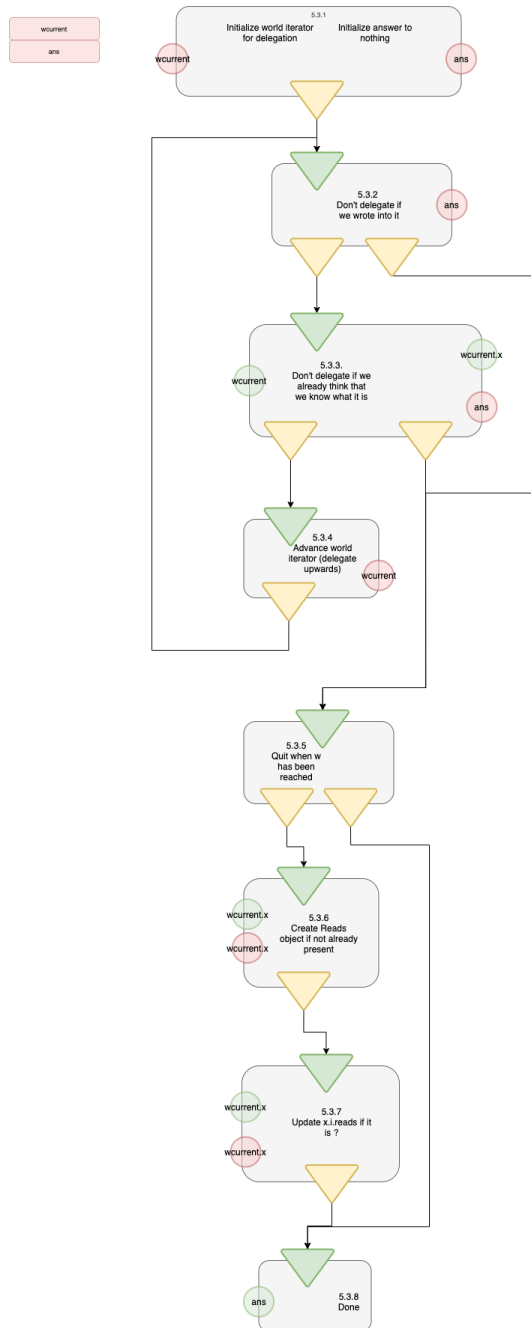
5.3 Slot Lookup

8-step algorithm, much like a flowchart

Components

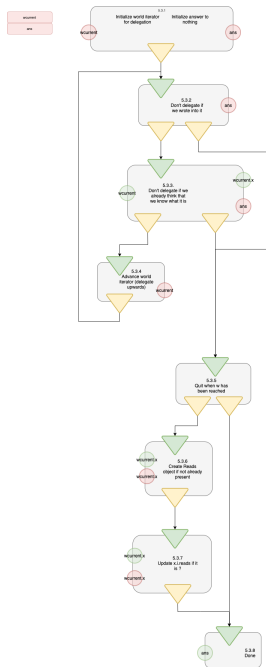


First Cut Control Flow




Control Flow (black boxes)

overview



Control Flow Detail



5.5 Commit

Algorithm

Commit 5.5

5.5.1:

either { child.reads.field == ?}
or { child.reads.field == parent.field }

5.5.2:

forall field { parent.writes.field := child.writes.field }

5.5.3:

forall field { parent.reads.field := child.reads.field }


** I don't yet understand the exception - is it necessary or merely an optimization?*

5.5.4:

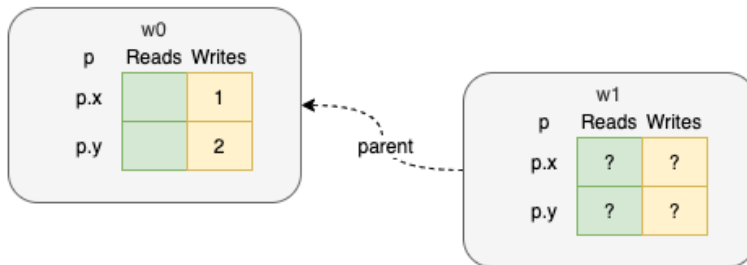
clear child

Section 9 says that *commit* in the top-level is a no-op.

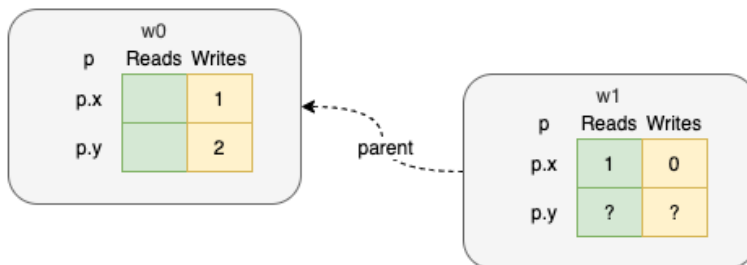
Successful Commit



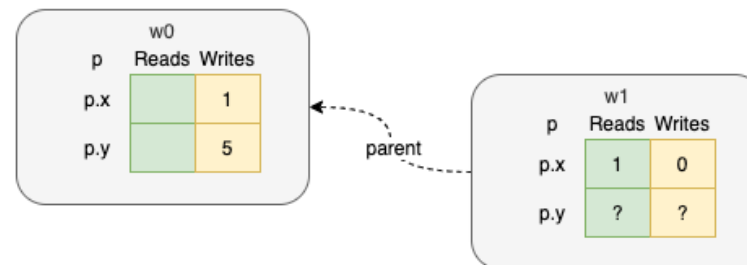
`w1 := w0.sprout()`



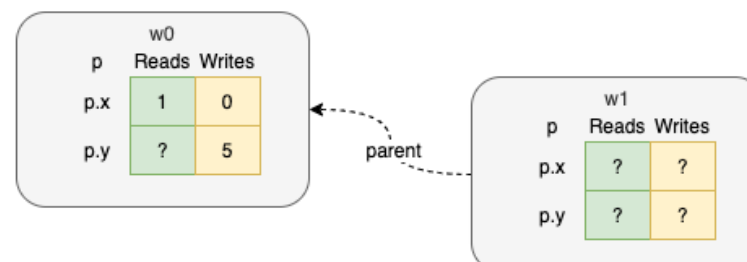
`in w2 { p.x := p.x - 1 }`




`in w0 { p.y := 5 }`




`w1.commit`




Failed Commit




w1 := w0.sprout()




in w2 { p.x := p.x + p.y }



in w0 { p.y := 5 }



w1.commit



Commit fails 5.5.1

6 Case Study #1 Bitmap Editor

6.3 Flattening Optimization

Fig. 6 shows 3 sub-worlds, where "1" is drawn in sub-world-1, "2" is drawn in sub-world-2 and "3" is drawn in sub-world-3.

In sub-world 3, the serifs are removed from the "1". This results in a few changes to the bitmap, (seen as white squares) but leaves much of the bitmap pixels in a "don't know" state.

The sub-world-3 bitmap is "flattened" to produce a Squeak-compatible bitmap and to improve editor performance.

7 Case Study #2 OMeta + Worlds

OMeta is a PEG parser technology (which inspired Ohm-JS).

PEG uses backtracking to parse incoming grammars, e.g. (A | B) tries to parse an A rule, if that fails, it backtracks and tries to parse a B rule.

Section 7 describes the experiment of replacing backtracking in OMeta with Worlds.

The conclusion:

And thus, with very little additional complexity, worlds can be use to make the use of side effects safer and easier to reason about in the presence of backtracking.

8 Related Work

Touches on:

- STM
- revisions and isolation types
- GL programming language (snapshots of the store)
- Contextual values
- Us, COP (Context-Oriented Programming)
- FDS (Functional Data Structures)
- concurrency control

9 Conclusion

- Worlds/JS, Worlds/Squeak
- hardware-assist for Worlds?
- infinite undo?
- persistence?
- in-memory vs. network side effects?

Previous Versions of the Paper

VPRI RN-2008-001

Contains an Introductory graphic that might help motivation.

http://www.vpri.org/pdf/rn2008001_worlds.pdf

Does The Concept Meet My Expectations?

Meet Expectations?

The Worlds concept partially meets my expectations and shows future promise.

Has the potential to subsume GC (Garbage Collection).

Scoped GC (ignoring commit).

Akin to UNIX® processes (which "clean up" when apps die). I like Worlds and UNIX® processes (for *isolation*), but, also, see potential for further improvement.

I see this as a basic technology that can be shaped (further scoped) to provide multiple notations (languages).

Commit of many variables has the potential to break locality-of-reference.

Is the implementation easier / quicker if *commit* is dropped?

Is *commit* needed?

JS Global and Window

Javascript nearly meets the needs of creating worlds, because all JS variables are contained in the *global* (node.js) object and the *window* (HTML) object.

For example

```
var x = 5;
```

is semantically equivalent to

```
var global.x = 5;
```

If we could change the value of `global`, we could create Worlds.

I *think* that the paper was based on the use of OMeta (PEG) to pre-process code and to transpile Worlds-based-code into stock JS.

Future

Essence of idea might lead to *isolation*, and isolated components.

Restrict coupling only to *ports*, not all variables.

Might be easy to build in, say, JavaScript'. In JS, all variables belong to a context, all variables are fields of a JS *object* { name:value, ... }

Thought: implement separate worlds but drop *commit*. Would this be enough to provide UNIX®-like *isolation*?

See Also

TXL (txl.ca) - functional language for parsing and rewriting syntax. TXL was intended for exploring new languages (by modifying existing languages). [*TXL was later used for Y2K detection, tree-rewriting, etc.*]